

davis\_xgb

May 16, 2024

```
[ ]: !pip install xgboost
```

```
[13]: import seaborn as sns
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import warnings
from sklearn.model_selection import train_test_split
import xgboost as xgb
```

## 1 Attempt #1 - no additional feature engineering

```
[15]: df = pd.read_csv('Davis.csv', parse_dates=['date'])
df.set_index('date', inplace=True)
#drop the columns that are not needed, unnamed:0 and hospital
df = df.drop(['Unnamed: 0', 'hospital'], axis=1)
```

```
[16]: df.dtypes
```

```
[16]: year                int64
monthday               int64
month                 int64
day                   int64
attendences           float64
min                   float64
max                   float64
aver                  float64
Hosp_ID               int64
Time_ID               int64
DAT                   float64
ThreeDAT              float64
EHIaccl               float64
dow                   int64
Sun                   int64
Mon                   int64
Tue                   int64
Wed                   int64
```

```

Thu          int64
Fri          int64
Sat          int64
Jan          int64
Feb          int64
Mar          int64
Apr          int64
May          int64
Jun          int64
Jul          int64
Aug          int64
Sep          int64
Oct          int64
Nov          int64
Dec          int64
Year_1       int64
Year_2       int64
Year_3       int64
Year_4       int64
Year_5       int64
Year_6       int64
Year_7       int64
Year_8       int64
dtype: object

```

```
[17]: #separate out the features and target variable
```

```
X, y = df.drop('attendences', axis=1), df[['attendences']]
```

```
[18]: #Split data into train/test split
```

```
X_train, X_test, y_train, y_test = X[:, '2014'], X[:, '2015'], y[:, '2014'], y[:, '2015']
↪
```

```
[19]: # Create regression matrices
```

```
dtrain_reg = xgb.DMatrix(X_train, y_train, enable_categorical=True)
dtest_reg = xgb.DMatrix(X_test, y_test, enable_categorical=True)
```

```
[23]: # Define hyperparameters
```

```
params = {"objective": "reg:squarederror", "tree_method": "hist"}

n = 100
model = xgb.train(
    params=params,
    dtrain=dtrain_reg,
    num_boost_round=n,
)
```

```
[24]: from sklearn.metrics import mean_squared_error
```

```
preds = model.predict(dtest_reg)
```

```
[25]: rmse = mean_squared_error(y_test, preds, squared=False)
```

```
print(f"RMSE of the base model: {rmse:.3f}")
```

RMSE of the base model: 21.290

c:\Users\kentm\Documents\Jupyter Notebooks\ed visit timeseries\.conda\Lib\site-packages\sklearn\metrics\\_regression.py:483: FutureWarning: 'squared' is deprecated in version 1.4 and will be removed in 1.6. To calculate the root mean squared error, use the function 'root\_mean\_squared\_error'.

warnings.warn(

```
[26]: evals = [(dtrain_reg, "train"), (dtest_reg, "validation")]
```

```
[27]: evals = [(dtrain_reg, "train"), (dtest_reg, "validation")]
```

```
model = xgb.train(  
    params=params,  
    dtrain=dtrain_reg,  
    num_boost_round=n,  
    evals=evals,  
)
```

|      |                     |                          |
|------|---------------------|--------------------------|
| [0]  | train-rmse:20.57698 | validation-rmse:35.38845 |
| [1]  | train-rmse:17.89705 | validation-rmse:31.04664 |
| [2]  | train-rmse:16.11103 | validation-rmse:28.23228 |
| [3]  | train-rmse:14.94763 | validation-rmse:26.22315 |
| [4]  | train-rmse:14.07073 | validation-rmse:23.80828 |
| [5]  | train-rmse:13.50216 | validation-rmse:23.01773 |
| [6]  | train-rmse:13.04559 | validation-rmse:22.37125 |
| [7]  | train-rmse:12.75375 | validation-rmse:22.09961 |
| [8]  | train-rmse:12.36247 | validation-rmse:21.64580 |
| [9]  | train-rmse:12.17684 | validation-rmse:21.39410 |
| [10] | train-rmse:11.91217 | validation-rmse:21.36551 |
| [11] | train-rmse:11.80576 | validation-rmse:21.07718 |
| [12] | train-rmse:11.53037 | validation-rmse:20.94231 |
| [13] | train-rmse:11.32419 | validation-rmse:20.94428 |
| [14] | train-rmse:11.16781 | validation-rmse:20.82353 |
| [15] | train-rmse:10.90405 | validation-rmse:21.12225 |
| [16] | train-rmse:10.69685 | validation-rmse:21.15110 |
| [17] | train-rmse:10.38318 | validation-rmse:21.24322 |
| [18] | train-rmse:10.26230 | validation-rmse:20.97955 |
| [19] | train-rmse:10.00594 | validation-rmse:20.88005 |
| [20] | train-rmse:9.94340  | validation-rmse:20.89322 |
| [21] | train-rmse:9.90405  | validation-rmse:20.91943 |

|      |                    |                          |
|------|--------------------|--------------------------|
| [22] | train-rmse:9.85032 | validation-rmse:20.59603 |
| [23] | train-rmse:9.81836 | validation-rmse:20.60465 |
| [24] | train-rmse:9.68540 | validation-rmse:20.49280 |
| [25] | train-rmse:9.42245 | validation-rmse:20.47224 |
| [26] | train-rmse:9.40610 | validation-rmse:20.46674 |
| [27] | train-rmse:9.31659 | validation-rmse:20.47315 |
| [28] | train-rmse:9.17599 | validation-rmse:20.45866 |
| [29] | train-rmse:9.02224 | validation-rmse:20.49002 |
| [30] | train-rmse:8.96467 | validation-rmse:20.79847 |
| [31] | train-rmse:8.91975 | validation-rmse:20.81242 |
| [32] | train-rmse:8.70874 | validation-rmse:20.81000 |
| [33] | train-rmse:8.59723 | validation-rmse:20.79737 |
| [34] | train-rmse:8.43756 | validation-rmse:20.83933 |
| [35] | train-rmse:8.36136 | validation-rmse:20.85988 |
| [36] | train-rmse:8.34698 | validation-rmse:20.87487 |
| [37] | train-rmse:8.32844 | validation-rmse:20.85109 |
| [38] | train-rmse:8.15004 | validation-rmse:20.85886 |
| [39] | train-rmse:8.04151 | validation-rmse:20.85571 |
| [40] | train-rmse:7.96766 | validation-rmse:20.88116 |
| [41] | train-rmse:7.83688 | validation-rmse:20.90447 |
| [42] | train-rmse:7.74672 | validation-rmse:20.91494 |
| [43] | train-rmse:7.67471 | validation-rmse:20.80084 |
| [44] | train-rmse:7.52446 | validation-rmse:20.82732 |
| [45] | train-rmse:7.46895 | validation-rmse:20.81068 |
| [46] | train-rmse:7.40479 | validation-rmse:20.81407 |
| [47] | train-rmse:7.27238 | validation-rmse:20.81378 |
| [48] | train-rmse:7.18302 | validation-rmse:20.83357 |
| [49] | train-rmse:7.08652 | validation-rmse:20.98890 |
| [50] | train-rmse:7.00587 | validation-rmse:21.01872 |
| [51] | train-rmse:6.92945 | validation-rmse:21.02215 |
| [52] | train-rmse:6.87135 | validation-rmse:21.03175 |
| [53] | train-rmse:6.73948 | validation-rmse:21.04085 |
| [54] | train-rmse:6.59420 | validation-rmse:21.07513 |
| [55] | train-rmse:6.51885 | validation-rmse:21.06905 |
| [56] | train-rmse:6.47816 | validation-rmse:21.07362 |
| [57] | train-rmse:6.36690 | validation-rmse:21.17923 |
| [58] | train-rmse:6.36098 | validation-rmse:21.18134 |
| [59] | train-rmse:6.22542 | validation-rmse:21.18277 |
| [60] | train-rmse:6.19454 | validation-rmse:21.20757 |
| [61] | train-rmse:6.09971 | validation-rmse:21.16663 |
| [62] | train-rmse:6.04597 | validation-rmse:21.17911 |
| [63] | train-rmse:5.92140 | validation-rmse:21.18289 |
| [64] | train-rmse:5.91505 | validation-rmse:21.18653 |
| [65] | train-rmse:5.85327 | validation-rmse:21.18679 |
| [66] | train-rmse:5.78370 | validation-rmse:21.18500 |
| [67] | train-rmse:5.70560 | validation-rmse:21.18758 |
| [68] | train-rmse:5.63868 | validation-rmse:21.18658 |
| [69] | train-rmse:5.54738 | validation-rmse:21.35656 |

|      |                    |                          |
|------|--------------------|--------------------------|
| [70] | train-rmse:5.43029 | validation-rmse:21.39749 |
| [71] | train-rmse:5.35898 | validation-rmse:21.39642 |
| [72] | train-rmse:5.23461 | validation-rmse:21.40526 |
| [73] | train-rmse:5.19625 | validation-rmse:21.40676 |
| [74] | train-rmse:5.15821 | validation-rmse:21.39932 |
| [75] | train-rmse:5.14638 | validation-rmse:21.33115 |
| [76] | train-rmse:5.13535 | validation-rmse:21.33361 |
| [77] | train-rmse:5.10735 | validation-rmse:21.34454 |
| [78] | train-rmse:5.08941 | validation-rmse:21.35428 |
| [79] | train-rmse:5.03167 | validation-rmse:21.30998 |
| [80] | train-rmse:4.94863 | validation-rmse:21.30703 |
| [81] | train-rmse:4.90116 | validation-rmse:21.30902 |
| [82] | train-rmse:4.82021 | validation-rmse:21.33733 |
| [83] | train-rmse:4.78846 | validation-rmse:21.32348 |
| [84] | train-rmse:4.72223 | validation-rmse:21.31695 |
| [85] | train-rmse:4.63709 | validation-rmse:21.31840 |
| [86] | train-rmse:4.54898 | validation-rmse:21.30781 |
| [87] | train-rmse:4.51128 | validation-rmse:21.40188 |
| [88] | train-rmse:4.44278 | validation-rmse:21.39443 |
| [89] | train-rmse:4.41445 | validation-rmse:21.40290 |
| [90] | train-rmse:4.38960 | validation-rmse:21.38219 |
| [91] | train-rmse:4.33664 | validation-rmse:21.37649 |
| [92] | train-rmse:4.29016 | validation-rmse:21.29750 |
| [93] | train-rmse:4.25880 | validation-rmse:21.29425 |
| [94] | train-rmse:4.16962 | validation-rmse:21.31368 |
| [95] | train-rmse:4.14030 | validation-rmse:21.31482 |
| [96] | train-rmse:4.06991 | validation-rmse:21.31194 |
| [97] | train-rmse:4.00219 | validation-rmse:21.32143 |
| [98] | train-rmse:3.94368 | validation-rmse:21.34343 |
| [99] | train-rmse:3.86686 | validation-rmse:21.28979 |

```
[35]: evals = [(dtrain_reg, "train"), (dtest_reg, "validation")]
n = 10000

model = xgb.train(
    params=params,
    dtrain=dtrain_reg,
    num_boost_round=n,
    evals=evals,
    verbose_eval=50,
    # Activate early stopping
    early_stopping_rounds=100
)
```

|       |                     |                          |
|-------|---------------------|--------------------------|
| [0]   | train-rmse:20.57698 | validation-rmse:35.38845 |
| [50]  | train-rmse:7.00587  | validation-rmse:21.01872 |
| [100] | train-rmse:3.82457  | validation-rmse:21.28313 |

```
[127]    train-rmse:2.90429    validation-rmse:21.40159
```

```
[31]: params = {"objective": "reg:squarederror", "tree_method": "hist"}
n = 1000

results = xgb.cv(
    params, dtrain_reg,
    num_boost_round=n,
    nfold=5,
    early_stopping_rounds=100
)
```


```
[34]: best_rmse = results['test-rmse-mean'].min()
best_rmse
```

```
[34]: 15.688509272020877
```

```
[39]: #apply model to the test data
dtest_reg = xgb.DMatrix(X_test)
preds_test = model.predict(dtest_reg)

#apply model to the train data
dtrain_reg = xgb.DMatrix(X_train)
preds_train = model.predict(dtrain_reg)
```

```
[43]: plt.figure(figsize=(20,10))
#graph preds and actuals
plt.plot(y_test.index, preds_test, label='Predictions')
plt.plot(y_test.index, y_test, label='Actuals')
plt.legend()
plt.title('Predictions vs Actuals on test set')
plt.show()
```



This graph is hard to look at... let's smooth it out a bit and add confidence intervals

```
[54]: # Calculate rolling averages
window_size = 14 # 7-day rolling window
y_test_smoothed = y_test.rolling(window=window_size).mean()
preds_test_smoothed = pd.Series(preds_test, index=y_test.index).
    ↪rolling(window=window_size).mean()


# Calculate residuals on training data
residuals = y_train['attendences'] - model.predict(xgb.DMatrix(X_train))

# Calculate the standard deviation of these residuals
error_std = np.std(residuals)

# Generate upper and lower confidence bounds
confidence_interval = 1.96 * error_std # 95% confidence interval
upper_bound = preds_test + confidence_interval
lower_bound = preds_test - confidence_interval
#smooth upper and lower bounds
upper_bound_smoothed = pd.Series(upper_bound, index=y_test.index).
    ↪rolling(window=window_size).mean()
lower_bound_smoothed = pd.Series(lower_bound, index=y_test.index).
    ↪rolling(window=window_size).mean()


plt.figure(figsize=(20,10))
plt.plot(y_test_smoothed.index, preds_test_smoothed, label='Predictions_
    ↪(Smoothed)', color='orange')
```

```
plt.plot(y_test_smoothed.index, y_test_smoothed, label='Actuals (Smoothed)',
        color='blue')
plt.fill_between(y_test.index, lower_bound_smoothed, upper_bound_smoothed,
                color='gray', alpha=0.3, label='95% Confidence Interval')
plt.legend()
plt.title('Smoothed Predicted ER Visits vs Actuals with 95% confidence
        interval')
plt.show()
```



```
[41]: #graph preds and actuals on train data
plt.figure(figsize=(20,10))
plt.plot(y_train.index, preds_train, label='Predictions', alpha=0.5)
plt.plot(y_train.index, y_train, label='Actuals', alpha=0.5)
plt.legend()
plt.title('Predictions vs Actuals on train set')
plt.show()
```





As expected, the predicted values for the training data line up amazingly well.

## 2 Attempt #2: adding lags and rolling averages

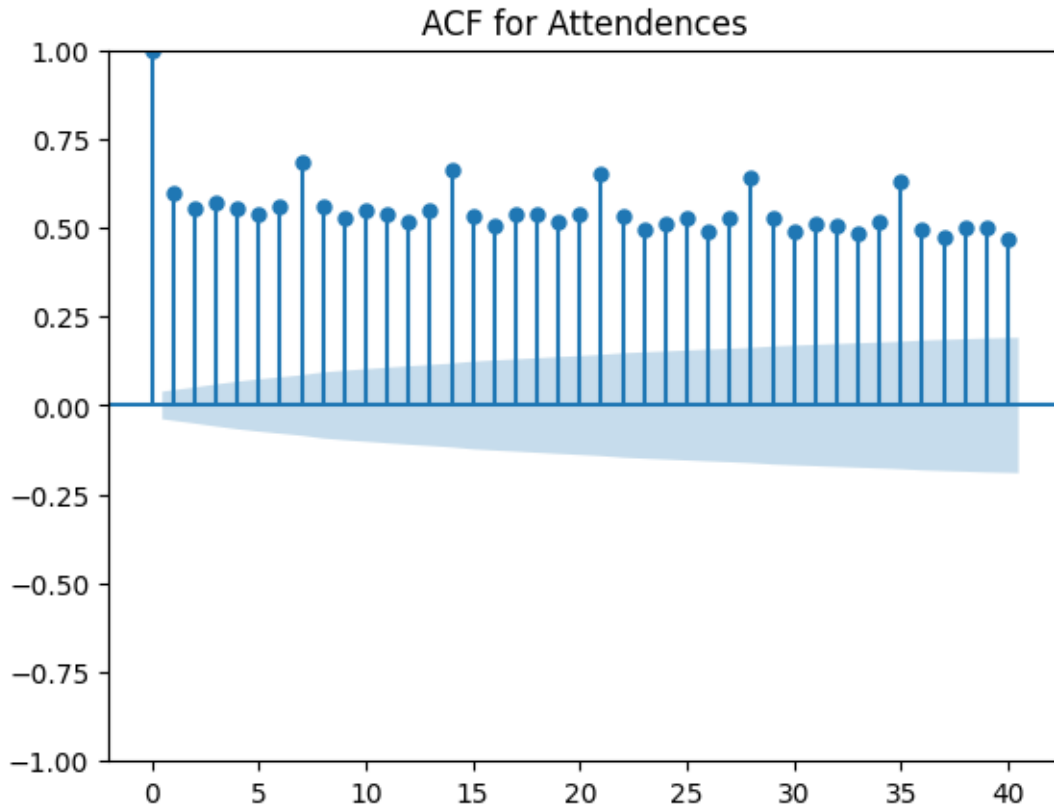
```
[79]: df = pd.read_csv('Davis.csv', parse_dates=['date'])
df.set_index('date', inplace=True)
#drop the columns that are not needed, unnamed:0 and hospital
df = df.drop(['Unnamed: 0', 'hospital'], axis=1)
```

Autocorrelation graphing to determine which lags are most relevant:

```
[77]: from statsmodels.graphics.tsaplots import plot_acf

plt.figure(figsize=(12, 6))
plot_acf(df['attendences'].dropna(), lags=40) # Adjust lags as necessary
plt.title('ACF for Attendences')
plt.show()
```

<Figure size 1200x600 with 0 Axes>



It looks like there is pretty heavy autocorrelation accross all 40, so I'll pick a few short term ones and weekly values to try to capture these trends without overcomplicating the model or leading to overfitting:

```
[102]: df = pd.read_csv('Davis.csv', parse_dates=['date'])
df.set_index('date', inplace=True)
#drop the columns that are not needed, unnamed:0 and hospital
df = df.drop(['Unnamed: 0', 'hospital'], axis=1)

#Creating lag features
df['lag_1'] = df['attendences'].shift(1)
df['lag_2'] = df['attendences'].shift(2)
df['lag_3'] = df['attendences'].shift(3)
df['lag_7'] = df['attendences'].shift(7)
df['lag_14'] = df['attendences'].shift(14)
df['lag_28'] = df['attendences'].shift(28)
df['lag_365'] = df['attendences'].shift(365)

#Adding somewhat arbitrary rolling averages
df['roll_avg_3'] = df['attendences'].rolling(window=3).mean()
#df['roll_avg_7'] = df['attendences'].rolling(window=7).mean()
```

```
#df['roll_avg_14'] = df['attendences'].rolling(window=14).mean()
#df['roll_avg_30'] = df['attendences'].rolling(window=30).mean()

#drop rows with any missing values in the lag or rolling avg columns
df = df.dropna()
```

```
[105]: df.head(5)
```

```
[105]:
```

|            | year | monthday | month | day | attendences | min | max  | aver | Hosp_ID | \ |
|------------|------|----------|-------|-----|-------------|-----|------|------|---------|---|
| date       |      |          |       |     |             |     |      |      |         |   |
| 2010-01-01 | 2010 | 101      | 1     | 1   | 156.0       | 7.0 | 13.0 | 11.0 | 6       |   |
| 2010-01-02 | 2010 | 102      | 1     | 2   | 168.0       | 7.0 | 14.0 | 11.0 | 6       |   |
| 2010-01-03 | 2010 | 103      | 1     | 3   | 168.0       | 3.0 | 10.0 | 7.0  | 6       |   |
| 2010-01-04 | 2010 | 104      | 1     | 4   | 171.0       | 3.0 | 12.0 | 7.0  | 6       |   |
| 2010-01-05 | 2010 | 105      | 1     | 5   | 165.0       | 4.0 | 8.0  | 6.0  | 6       |   |

|            | Time_ID | ... | Year_7 | Year_8 | lag_1 | lag_2 | lag_3 | lag_7 | lag_14 | \ |
|------------|---------|-----|--------|--------|-------|-------|-------|-------|--------|---|
| date       |         | ... |        |        |       |       |       |       |        |   |
| 2010-01-01 | 366     | ... | 0      | 0      | 134.0 | 157.0 | 142.0 | 128.0 | 181.0  |   |
| 2010-01-02 | 367     | ... | 0      | 0      | 156.0 | 134.0 | 157.0 | 144.0 | 145.0  |   |
| 2010-01-03 | 368     | ... | 0      | 0      | 168.0 | 156.0 | 134.0 | 146.0 | 131.0  |   |
| 2010-01-04 | 369     | ... | 0      | 0      | 168.0 | 168.0 | 156.0 | 178.0 | 173.0  |   |
| 2010-01-05 | 370     | ... | 0      | 0      | 171.0 | 168.0 | 168.0 | 142.0 | 133.0  |   |

|            | lag_28 | lag_365 | roll_avg_3 |
|------------|--------|---------|------------|
| date       |        |         |            |
| 2010-01-01 | 150.0  | 180.0   | 149.000000 |
| 2010-01-02 | 151.0  | 193.0   | 152.666667 |
| 2010-01-03 | 142.0  | 171.0   | 164.000000 |
| 2010-01-04 | 143.0  | 151.0   | 169.000000 |
| 2010-01-05 | 156.0  | 177.0   | 168.000000 |

[5 rows x 49 columns]

```
[103]: #separate out the features and target variable
X, y = df.drop('attendences', axis=1), df[['attendences']]
#Split data into train/test split
X_train, X_test, y_train, y_test = X[:'2014'], X['2015':], y[:'2014'], y['2015':]
# Create regression matrices
dtrain_reg = xgb.DMatrix(X_train, y_train, enable_categorical=True)
dtest_reg = xgb.DMatrix(X_test, y_test, enable_categorical=True)
#train model
evals = [(dtrain_reg, "train"), (dtest_reg, "validation")]
n = 10000
model = xgb.train(
    params=params,
```

```

dtrain=dtrain_reg,
num_boost_round=n,
evals=evals,
verbose_eval=50,
# Activate early stopping
early_stopping_rounds=5
)

```

```

[0]      train-rmse:19.17805      validation-rmse:31.01858
[34]      train-rmse:1.37975      validation-rmse:11.39988

```

```

[104]: #apply model to the test data
dtest_reg = xgb.DMatrix(X_test)
preds_test = model.predict(dtest_reg)

# Calculate rolling averages
window_size = 14 # 7-day rolling window
y_test_smoothed = y_test.rolling(window=window_size).mean()
preds_test_smoothed = pd.Series(preds_test, index=y_test.index).
    ↪rolling(window=window_size).mean()

# Calculate residuals on training data
residuals = y_train['attendences'] - model.predict(xgb.DMatrix(X_train))


# Calculate the standard deviation of these residuals
error_std = np.std(residuals)

# Generate upper and lower confidence bounds
confidence_interval = 1.96 * error_std # 95% confidence interval
upper_bound = preds_test + confidence_interval
lower_bound = preds_test - confidence_interval
#smooth upper and lower bounds
upper_bound_smoothed = pd.Series(upper_bound, index=y_test.index).
    ↪rolling(window=window_size).mean()
lower_bound_smoothed = pd.Series(lower_bound, index=y_test.index).
    ↪rolling(window=window_size).mean()

plt.figure(figsize=(20,10))
plt.plot(y_test_smoothed.index, preds_test_smoothed, label='Predictions_
    ↪(Smoothed)', color='orange')
plt.plot(y_test_smoothed.index, y_test_smoothed, label='Actuals (Smoothed)',
    ↪color='blue')
plt.fill_between(y_test.index, lower_bound_smoothed, upper_bound_smoothed,
    ↪color='gray', alpha=0.3, label='95% Confidence Interval')
plt.legend()
plt.title('Smoothed Predicted ER Visits vs Actuals with 95% confidence_
    ↪interval')

```

```
plt.show()
```



### 3 Attempt #3: Using above augmented data with grid search hyperparameter optimization

```
[106]: df = pd.read_csv('Davis.csv', parse_dates=['date'])
df.set_index('date', inplace=True)
#drop the columns that are not needed, unnamed:0 and hospital
df = df.drop(['Unnamed: 0', 'hospital'], axis=1)

#Creating lag features
df['lag_1'] = df['attendences'].shift(1)
df['lag_2'] = df['attendences'].shift(2)
df['lag_3'] = df['attendences'].shift(3)
df['lag_7'] = df['attendences'].shift(7)
df['lag_14'] = df['attendences'].shift(14)
df['lag_28'] = df['attendences'].shift(28)
df['lag_365'] = df['attendences'].shift(365)

#Adding somewhat arbitrary rolling averages
df['roll_avg_3'] = df['attendences'].rolling(window=3).mean()
#df['roll_avg_7'] = df['attendences'].rolling(window=7).mean()
#df['roll_avg_14'] = df['attendences'].rolling(window=14).mean()
#df['roll_avg_30'] = df['attendences'].rolling(window=30).mean()

#drop rows with any missing values in the lag or rolling avg columns
```

```
df = df.dropna()
```

```
[107]: from sklearn.model_selection import GridSearchCV
from xgboost import XGBRegressor

#separate out the features and target variable
X, y = df.drop('attendences', axis=1), df[['attendences']]
#Split data into train/test split
X_train, X_test, y_train, y_test = X[:'2014'], X['2015':], y[:'2014'], y['2015':
↪]
# Create regression matrices
dtrain_reg = xgb.DMatrix(X_train, y_train, enable_categorical=True)
dtest_reg = xgb.DMatrix(X_test, y_test, enable_categorical=True)

# Define the model
model = XGBRegressor()

# Define the parameter grid
param_grid = {
    'max_depth': [3, 4, 5],
    'learning_rate': [0.01, 0.1, 0.2],
    'n_estimators': [100, 200],
    'subsample': [0.7, 0.8, 0.9]
}

# Setup the grid search
grid_search = GridSearchCV(estimator=model, param_grid=param_grid, cv=3,
↪scoring='neg_mean_squared_error', verbose=1)

# Fit the grid search to the data
grid_search.fit(X_train, y_train)

# Get the best parameters
print("Best parameters:", grid_search.best_params_)
```

Fitting 3 folds for each of 54 candidates, totalling 162 fits

Best parameters: {'learning\_rate': 0.1, 'max\_depth': 4, 'n\_estimators': 200, 'subsample': 0.8}

```
[114]: #train model
evals = [(dtrain_reg, "train"), (dtest_reg, "validation")]
n = 10000
model = xgb.train(
    params=grid_search.best_params_,
    dtrain=dtrain_reg,
    num_boost_round=n,
    evals=evals,
```

```

verbose_eval=50,
# Activate early stopping
early_stopping_rounds=25
)

```

```

[0]      train-rmse:23.26391      validation-rmse:36.72013
[50]      train-rmse:6.04741      validation-rmse:12.85865
[100]     train-rmse:3.43751      validation-rmse:11.24115

```

c:\Users\kentm\Documents\Jupyter Notebooks\ed visit timeseries\.conda\Lib\site-packages\xgboost\core.py:160: UserWarning: [11:11:55] WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0b3782d1791676daf-1\xgboost\xgboost-ci-windows\src\learner.cc:742: Parameters: { "n\_estimators" } are not used.

```
warnings.warn(smsg, UserWarning)
```

```

[150]     train-rmse:2.20958      validation-rmse:10.81809
[200]     train-rmse:1.60785      validation-rmse:10.53766
[250]     train-rmse:1.28960      validation-rmse:10.37736
[300]     train-rmse:1.07537      validation-rmse:10.32707
[349]     train-rmse:0.92537      validation-rmse:10.30443

```

```

[115]: #apply model to the test data
dtest_reg = xgb.DMatrix(X_test)
preds_test = model.predict(dtest_reg)

# Calculate rolling averages
window_size = 14 # 7-day rolling window
y_test_smoothed = y_test.rolling(window=window_size).mean()
preds_test_smoothed = pd.Series(preds_test, index=y_test.index).
    ↪rolling(window=window_size).mean()

# Calculate residuals on training data
residuals = y_train['attendences'] - model.predict(xgb.DMatrix(X_train))

# Calculate the standard deviation of these residuals
error_std = np.std(residuals)


# Generate upper and lower confidence bounds
confidence_interval = 1.96 * error_std # 95% confidence interval
upper_bound = preds_test + confidence_interval
lower_bound = preds_test - confidence_interval
#smooth upper and lower bounds
upper_bound_smoothed = pd.Series(upper_bound, index=y_test.index).
    ↪rolling(window=window_size).mean()
lower_bound_smoothed = pd.Series(lower_bound, index=y_test.index).
    ↪rolling(window=window_size).mean()

```

```

plt.figure(figsize=(20,10))
plt.plot(y_test_smoothed.index, preds_test_smoothed, label='Predictions_
↳(Smoothed)', color='orange')
plt.plot(y_test_smoothed.index, y_test_smoothed, label='Actuals (Smoothed)',
↳color='blue')
plt.fill_between(y_test.index, lower_bound_smoothed, upper_bound_smoothed,
↳color='gray', alpha=0.3, label='95% Confidence Interval')
plt.legend()
plt.title('Smoothed Predicted ER Visits vs Actuals with 95% confidence_
↳interval')
plt.show()

```



#### 4 Attempt #4: Cross validating performance accross multiple time series splits

```

[129]: import pandas as pd
import numpy as np
from sklearn.model_selection import TimeSeriesSplit
import xgboost as xgb
from sklearn.metrics import mean_squared_error
from xgboost import XGBRegressor
from sklearn.model_selection import GridSearchCV

# Load and prepare the data
df = pd.read_csv('Davis.csv', parse_dates=['date'])

```



```

df.set_index('date', inplace=True)
df = df.drop(['Unnamed: 0', 'hospital'], axis=1)

# Add lag and rolling features to the whole dataset
lags = [1, 2, 3, 7, 14, 28, 365]
for lag in lags:
    df[f'lag_{lag}'] = df['attendences'].shift(lag)

rolling_windows = [3]
for window in rolling_windows:
    df[f'rolling_avg_{window}'] = df['attendences'].rolling(window=window).
    ↪mean()

# Drop rows with NaN values that were created by shift and rolling
df.dropna(inplace=True)

# Time Series Cross-validation
tscv = TimeSeriesSplit(n_splits=5)

for train_index, test_index in tscv.split(df):
    train, test = df.iloc[train_index], df.iloc[test_index]
    X_train, y_train = train.drop('attendences', axis=1), train['attendences']
    X_test, y_test = test.drop('attendences', axis=1), test['attendences']

    # Define the model
    model = XGBRegressor()

    # Parameter grid
    param_grid = {
        'max_depth': [3, 4, 5],
        'learning_rate': [0.01, 0.1, 0.2],
        'n_estimators': [100, 200],
        'subsample': [0.7, 0.8, 0.9]
    }

    # Grid search
    grid_search = GridSearchCV(estimator=model, param_grid=param_grid, cv=3,
    ↪scoring='neg_mean_squared_error', verbose=1)
    grid_search.fit(X_train, y_train)
    print("Best parameters:", grid_search.best_params_)

    # Use best parameters to train the model
    best_params = grid_search.best_params_
    final_model = XGBRegressor(**best_params)
    final_model.fit(X_train, y_train)

    # Predict and evaluate

```

```

predictions = final_model.predict(X_test)
rmse = np.sqrt(mean_squared_error(y_test, predictions))
print(f"Fold RMSE: {rmse:.3f}")

```

Fitting 3 folds for each of 54 candidates, totalling 162 fits  
 Best parameters: {'learning\_rate': 0.1, 'max\_depth': 3, 'n\_estimators': 200, 'subsample': 0.7}  
 Fold RMSE: 11.868  
 Fitting 3 folds for each of 54 candidates, totalling 162 fits  
 Best parameters: {'learning\_rate': 0.2, 'max\_depth': 3, 'n\_estimators': 200, 'subsample': 0.7}  
 Fold RMSE: 6.593  
 Fitting 3 folds for each of 54 candidates, totalling 162 fits  
 Best parameters: {'learning\_rate': 0.2, 'max\_depth': 3, 'n\_estimators': 200, 'subsample': 0.9}  
 Fold RMSE: 6.141  
 Fitting 3 folds for each of 54 candidates, totalling 162 fits  
 Best parameters: {'learning\_rate': 0.2, 'max\_depth': 3, 'n\_estimators': 200, 'subsample': 0.7}  
 Fold RMSE: 10.375  
 Fitting 3 folds for each of 54 candidates, totalling 162 fits  
 Best parameters: {'learning\_rate': 0.2, 'max\_depth': 3, 'n\_estimators': 200, 'subsample': 0.8}  
 Fold RMSE: 5.128

## 4.1 Training final model

```

[132]: # Assuming best_params are defined from your previous GridSearch
df = pd.read_csv('Davis.csv', parse_dates=['date'])
df.set_index('date', inplace=True)
#drop the columns that are not needed, unnamed:0 and hospital
df = df.drop(['Unnamed: 0', 'hospital'], axis=1)

#Creating lag features
df['lag_1'] = df['attendences'].shift(1)
df['lag_2'] = df['attendences'].shift(2)
df['lag_3'] = df['attendences'].shift(3)
df['lag_7'] = df['attendences'].shift(7)
df['lag_14'] = df['attendences'].shift(14)
df['lag_28'] = df['attendences'].shift(28)
df['lag_365'] = df['attendences'].shift(365)

#Adding somewhat arbitrary rolling averages
df['roll_avg_3'] = df['attendences'].rolling(window=3).mean()
#df['roll_avg_7'] = df['attendences'].rolling(window=7).mean()
#df['roll_avg_14'] = df['attendences'].rolling(window=14).mean()
#df['roll_avg_30'] = df['attendences'].rolling(window=30).mean()

```

```

#drop rows with any missing values in the lag or rolling avg columns
df = df.dropna()

#separate out the features and target variable
X, y = df.drop('attendences', axis=1), df[['attendences']]
#Split data into train/test split
X_train, X_test, y_train, y_test = X[:'2014'], X['2015':], y[:'2014'], y['2015':
↪]
# Create regression matrices
dtrain_reg = xgb.DMatrix(X_train, y_train, enable_categorical=True)
dtest_reg = xgb.DMatrix(X_test, y_test, enable_categorical=True)

best_params = {'learning_rate': 0.2, 'max_depth': 3, 'n_estimators': 200,
↪ 'subsample': 0.7}
final_model = XGBRegressor(**best_params)
final_model.fit(X_train, y_train)

predictions = final_model.predict(X_test)
rmse = np.sqrt(mean_squared_error(y_test, predictions))
print(f"Fold RMSE: {rmse:.3f}")

```

Fold RMSE: 10.965

```

[137]: #apply model to the test data
dtest_reg = xgb.DMatrix(X_test)
preds_test = final_model.predict(X_test)

# Calculate rolling averages
window_size = 14 # 7-day rolling window
y_test_smoothed = y_test.rolling(window=window_size).mean()
preds_test_smoothed = pd.Series(preds_test, index=y_test.index).
↪rolling(window=window_size).mean()

# Calculate residuals on training data
residuals = y_train['attendences'] - final_model.predict(X_train)

# Calculate the standard deviation of these residuals
error_std = np.std(residuals)

# Generate upper and lower confidence bounds
confidence_interval = 1.96 * error_std # 95% confidence interval
upper_bound = preds_test + confidence_interval
lower_bound = preds_test - confidence_interval
#smooth upper and lower bounds
upper_bound_smoothed = pd.Series(upper_bound, index=y_test.index).
↪rolling(window=window_size).mean()

```

```

lower_bound_smoothed = pd.Series(lower_bound, index=y_test.index).
    ↪rolling(window=window_size).mean()

plt.figure(figsize=(20,10))
plt.plot(y_test_smoothed.index, preds_test_smoothed, label='Predictions_
    ↪(Smoothed)', color='orange')
plt.plot(y_test_smoothed.index, y_test_smoothed, label='Actuals (Smoothed)',
    ↪color='blue')
plt.fill_between(y_test.index, lower_bound_smoothed, upper_bound_smoothed,
    ↪color='gray', alpha=0.3, label='95% Confidence Interval')
plt.legend()
plt.title('Smoothed Predicted ER Visits vs Actuals with 95% confidence_
    ↪interval')
plt.show()

```

